

In [10]:

```
# Estimation of Change in Price of Rice Commodity in India

# Importing all libraries used
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import pyplot
from datetime import datetime
from pandas import Series
from pandas import DataFrame
from pandas import concat
import seaborn as sns
from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import Normalizer
from pandas import read_csv
from numpy import set_printoptions
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.svm import LinearSVC
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.linear_model import LogisticRegression
from tabulate import tabulate
%matplotlib inline

pd.set_option('display.max_rows', None)
```

In [11]:

```
# Load food-prices dataset
df = pd.read_csv("wfp_market_food_prices.csv", encoding="latin-1")

# Drop unnecessary features
df.drop(['mp_commoditysource', 'cm_id', 'mkt_id', 'mkt_name', 'cur_id', 'pt_id', 'pt_name', 'um_id', 'um_name', 'adm0_id', 'adm1_id', 'adm1_name', 'cur_name'], inplace=True, axis=1)

# Consider only Rice commodity in India
df = df[df.adm0_name == "India"]
df = df[df.cm_name == 'Rice']
```

In [12]:

```
# Preprocessing of data
```

```
df['date'] = df.apply(lambda x: datetime(x['mp_year'],x["mp_month"],1),axis=1)
df = df.rename(columns={'mp_month':'month','mp_year':'year','mp_price':'price'})
df.drop(['adm0_name', 'cm_name'],inplace=True,axis=1)
df.index = np.arange(1, len(df) + 1)
```

In [15]:

```
# Feature Engineering
```

```
# Read carbon dioxide data
```

```
df_co2 = pd.read_csv("carbon_india.csv", sep=';')
```

```
# Merge co2 values into main dataframe
```

```
df = df.merge(df_co2.set_index('year'), on='year', how='left')
```

```
# Read temperature data
```

```
df_temp = pd.read_excel("temperature.xls")
```

```
df_temp = df_temp.rename(columns={' Month': 'month','\tYear':'year','tas':'temperature'})
```

```
df_temp['date'] = df_temp.apply(lambda x: datetime(int(x['year']),int(x['month']),1),axis=1)
```

```
df_temp.drop([' Country', 'month', 'year'],axis=1,inplace=True)
```

```
df = df.merge(df_temp.set_index('date'), on='date', how='left')
```

```
# Read inflation rate data
```

```
df_inf = pd.read_excel("inflation_rate_india.xls")
```

```
df = df.merge(df_inf.set_index('year'), on='year', how='left')
```

```
# Read consumer price index data
```

```
df_cpi = pd.read_excel("consumer_price_index_india.xls")
```

```
df = df.merge(df_cpi.set_index('year'), on='year', how='left')
```

```
# Read rainfall data
```

```
df_rain = pd.read_excel("rainfall.xls")
```

```
df_rain = df_rain.rename(columns={' Month': 'month','\tYear':'year','pr':'rainfall'})
```


```
df_rain['date'] = df_rain.apply(lambda x: datetime(int(x['year']),int(x['month']),1),axis=1)
```

```
df_rain.drop([' Country', 'month', 'year'],axis=1,inplace=True)
```

```
df = df.merge(df_rain.set_index('date'), on='date', how='left')
```

In [16]:

```
# Plot price versus date
dfp = pd.DataFrame(pd.Series(data=df.loc[:, 'price'].values).values, columns=[ 'p
rice' ], index=df[ 'date' ])
dfp.plot()
plt.title("Price of Rice vs date in India")
plt.show()
```



In [17]:

```
# Calculate 3 month mean of price
width = 3
lag1 = dfp.shift(1)
lag3 = dfp.shift(width - 1)
window = lag1.rolling(window=width)
df3 = concat([window.mean()], axis=1)
df3.columns = ['3 month mean price']
df['3 month mean'] = df3.values

# Calculate 6 month mean of price
width = 6
lag1 = dfp.shift(1)
lag3 = dfp.shift(width - 1)
window = lag1.rolling(window = width)
df6 = concat([window.mean()], axis=1)
df6.columns = ['6 month mean price']
df['6 month mean'] = df6.values

# Calculate 9 month mean of price
width = 9
lag1 = dfp.shift(1)
lag3 = dfp.shift(width - 1)
window = lag1.rolling(window=width)
df9 = concat([window.mean()], axis=1)
df9.columns = ['9 month mean price']
df['9 month mean'] = df9.values

# Calculate 12 month mean of price
width = 12
lag1 = dfp.shift(1)
lag3 = dfp.shift(width - 1)
window = lag1.rolling(window=width)
df12 = concat([window.mean()], axis=1)
df12.columns = ['12 month mean price']
df['12 month mean'] = df12.values

df.dropna(axis=0,inplace=True)
```

In [9]:

```
# Description of dataset

actual_df = df
print('Shape:',df.shape)
print('Columns:',df.columns)
print('Datatypes:',df.dtypes)
display(df.describe(include="all").T)
```

Shape: (7369, 13)

Columns: Index(['month', 'year', 'price', 'date', 'co2', 'temperature', 'inflation', 'consumer_price_index', 'rainfall', '3 month mean', '6 month mean', '9 month mean', '12 month mean'], dtype='object')

Datatypes: month int64

year int64

price float64

date datetime64[ns]

co2 float64

temperature float64

inflation float64

consumer_price_index float64

rainfall float64

3 month mean float64

6 month mean float64

9 month mean float64

12 month mean float64

dtype: object

	count	unique	top	freq	first	last	mean	
month	7369	NaN	NaN	NaN	NaN	NaN	6.34957	3.
year	7369	NaN	NaN	NaN	NaN	NaN	2008.87	6.
price	7369	NaN	NaN	NaN	NaN	NaN	19.6229	8.
date	7369	275	2017-01-01 00:00:00	59	1994-01-01 00:00:00	2017-06-01 00:00:00	NaN	N.
co2	7369	NaN	NaN	NaN	NaN	NaN	1744.13	5.
temperature	7369	NaN	NaN	NaN	NaN	NaN	25.4841	4.
inflation	7369	NaN	NaN	NaN	NaN	NaN	7.15222	2.
consumer_price_index	7369	NaN	NaN	NaN	NaN	NaN	104.009	40.
rainfall	7369	NaN	NaN	NaN	NaN	NaN	85.1589	8.
3 month mean	7369	NaN	NaN	NaN	NaN	NaN	19.6173	8.
6 month mean	7369	NaN	NaN	NaN	NaN	NaN	19.6133	8.
9 month mean	7369	NaN	NaN	NaN	NaN	NaN	19.6094	8.
12 month mean	7369	NaN	NaN	NaN	NaN	NaN	19.6056	8.

In [11]:

```
# Plot histogram of dataset
```


```
fig = plt.figure(figsize = (15,20))
ax = fig.gca()
df.hist(ax=ax)
```

```
D:\Anaconda\lib\site-packages\IPython\core\interactiveshell.py:296
1: UserWarning: To output multiple subplots, the figure containing
the passed axes is being cleared
    exec(code_obj, self.user_global_ns, self.user_ns)
```


Out[11]:

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC0A35BA8>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC1080FD0>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC10B06A0>],
       [<matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC10D6D30>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC110A400>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC110A438>],
       [<matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC132C940>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC1353FD0>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC13846A0>],
       [<matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC13ADD30>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC13DD400>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x0000012
DC1405A90>]],
      dtype=object)
```


12 month mean




3 month mean




6 month mean




9 month mean




co2




consumer_price_index




inflation




month




price




rainfall



temperature



year



In [6]:

```
# Calculate qcut which will be a label to indicate the percentage change in price
```

```
qcut = pd.qcut((df.price.sort_values(ascending=True).values), 10, labels=[1,2,3,4,5,6,7,8,9,10])
print(qcut)
df.sort_values(by=['price'], inplace=True)
df['label'] = qcut
df.sort_index(inplace=True)
df.dropna(axis=0, inplace=True)
```

```
[1, 1, 1, 1, 1, ..., 10, 10, 10, 10, 10]
```

```
Length: 7381
```

```
Categories (10, int64): [1 < 2 < 3 < 4 ... 7 < 8 < 9 < 10]
```

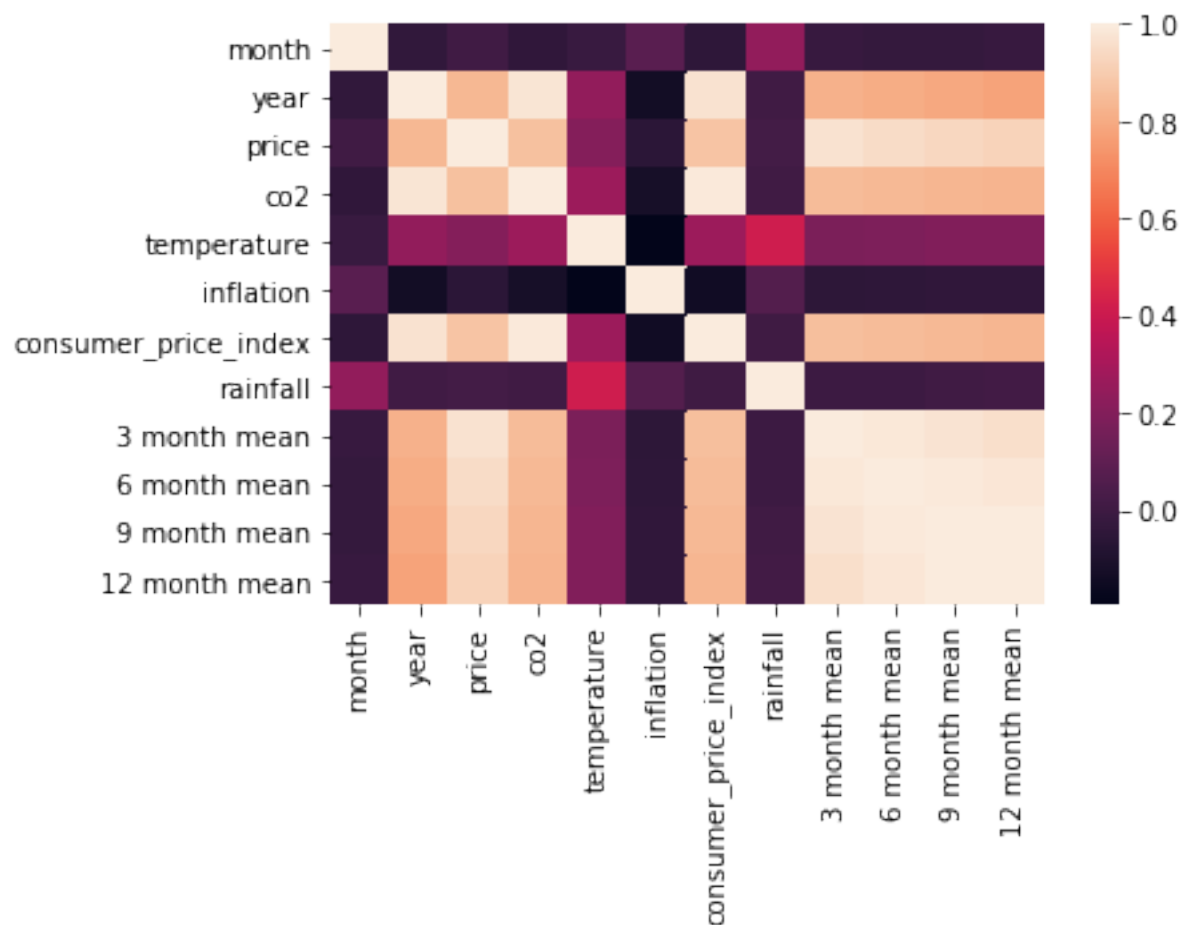
In [13]:

```
# Plot correlation matrix plot
```

```
correlations = df.corr()
sns.heatmap(correlations, xticklabels=correlations.columns.values, yticklabels=correlations.columns.values)
```

Out[13]:

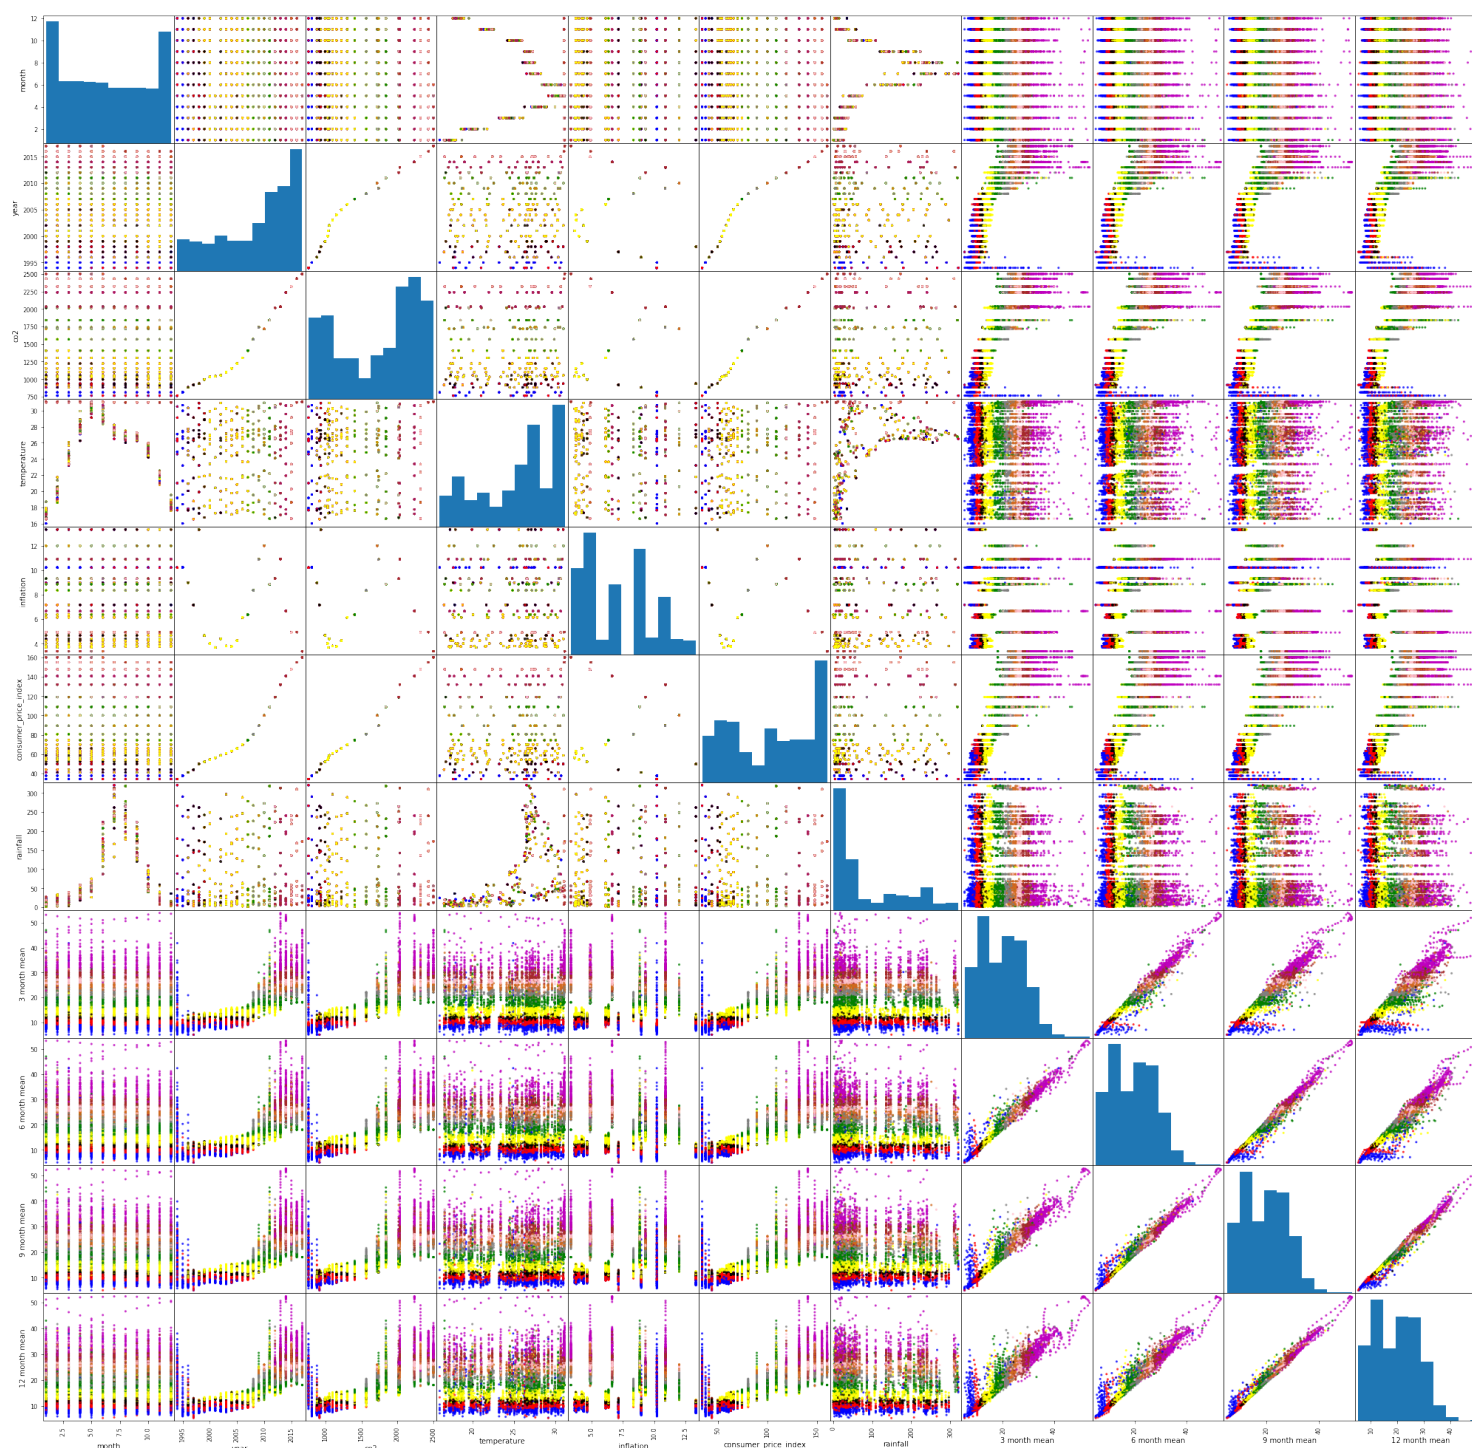
<matplotlib.axes._subplots.AxesSubplot at 0x12dc16650b8>




```
In [14]:
```

```
# Plot scatter matrix of the dataset
```

```
df.drop(columns=['date','price'],inplace=True)
palette = {1 : 'blue', 2 : 'red',3:'black',4:'yellow',5:'green',6:'gray',7:'chocolate',8:'pink',9:'brown',10:'m'}
labels=df['label']
labels_c = list(map(lambda x: palette[int(x)], labels))
grr = pd.plotting.scatter_matrix(df,figsize=(35,35), alpha=0.8, c=labels_c)
```



In [16]:

```
#Standardizing data for Gaussian distribution
```

```
array = df.values
# separate array into input and output components
cols = 11
X = array[:, :cols]
Y = array[:, cols]
scaler = StandardScaler().fit(X)
X = scaler.transform(X)
```

```
# summarize transformed data
np.set_printoptions(precision=3)
print(X[0:5, :])
```

```
# split train and test data
test_size = 0.33
seed = 7
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=test_size,
random_state=seed)
Y_train = Y_train.astype('int')
Y_test = Y_test.astype('int')
actual = Y_test
```

```
[[[-1.26   -1.786  -1.48   -1.483   0.004  -1.471  -0.906  -1.344  -1.317  -
1.326
   -1.345]
  [-0.97   -1.786  -1.48   -0.337   0.004  -1.471  -0.772  -1.364  -1.327  -
1.329
   -1.34   ]
  [-0.681  -1.786  -1.48    0.169   0.004  -1.471  -0.523  -1.344  -1.337  -
1.329
   -1.34   ]
  [-0.391  -1.786  -1.48    0.803   0.004  -1.471  -0.491  -1.304  -1.337  -
1.329
   -1.338]
  [-0.101  -1.786  -1.48    0.799   0.004  -1.471   0.715  -1.285  -1.337  -
1.329
   -1.335]]
```

D:\Anaconda\lib\site-packages\sklearn\utils\validation.py:475: DataConversionWarning: Data with input dtype object was converted to float64 by StandardScaler.

```
warnings.warn(msg, DataConversionWarning)
```

D:\Anaconda\lib\site-packages\sklearn\utils\validation.py:475: DataConversionWarning: Data with input dtype object was converted to float64 by StandardScaler.

```
warnings.warn(msg, DataConversionWarning)
```

In [17]:

```
# Gaussian NB model

Y = Y.astype('int')
kfold = KFold(n_splits=10, random_state=7)
model = GaussianNB()
model.fit(X_train,Y_train)
results = cross_val_score(model, X, Y, cv=kfold, scoring='accuracy')
print('Gaussian NB model Accuracy:\nMean=',round(results.mean()*100,3), '%\tStd deviation=',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actual,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : support}
print("\\n",tabulate(results,headers="keys"))
```

Gaussian NB model Accuracy:
Mean= 62.952 % Std deviation= 5.199 %

precision	recall	fscore	support
0.849057	0.54	0.660147	250
0.674129	0.818731	0.739427	331
0.528926	0.825806	0.644836	155
0.746914	0.461832	0.570755	262
0.625731	0.730375	0.674016	293
0.559585	0.404494	0.469565	267
0.484429	0.608696	0.539499	230
0.450292	0.442529	0.446377	174
0.633758	0.802419	0.708185	248
0.873418	0.621622	0.726316	222

In [18]:

```
# Decision Tree Classifier model

clf = DecisionTreeClassifier()
clf.fit(X_train, Y_train)
kfold = KFold(n_splits=10, random_state=7)
model = DecisionTreeClassifier()
model.fit(X_train,Y_train)
results = cross_val_score(model, X, Y, cv=kfold, scoring='accuracy')
print('Decision Tree Classifier Accuracy:\nMean=',round(results.mean()*100,3),
'%\tStd deviation= ',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actu
al,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : s
upport}
print("\\n",tabulate(results,headers="keys"))
```

Decision Tree Classifier Accuracy:

Mean= 68.354 % Std deviation= 3.637 %

precision	recall	fscore	support
0.818584	0.74	0.777311	250
0.727794	0.767372	0.747059	331
0.616438	0.580645	0.598007	155
0.758865	0.816794	0.786765	262
0.785714	0.788396	0.787053	293
0.707317	0.651685	0.678363	267
0.620087	0.617391	0.618736	230
0.613402	0.683908	0.646739	174
0.726908	0.729839	0.72837	248
0.820276	0.801802	0.810934	222

In [19]:

```
# SVC model

alphas = np.array([1,0.1,0.01,0.001,0.0001,0])
param_grid = [
    {'C': [1, 10, 100, 1000], 'kernel': ['linear']},
    {'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001], 'kernel': ['rbf']},
]
model = SVC()
grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid.fit(X_train, Y_train)
print(grid.best_params_)
print('SVC model Accuracy:',grid.best_score_*100)

prediction = grid.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actu
al,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : s
upport}
print("\\n",tabulate(results,headers="keys"))
```

```
{'C': 1000, 'gamma': 0.001, 'kernel': 'rbf'}
SVC model Accuracy: 76.88879886570791
```

precision	recall	fscore	support
0.88835	0.732	0.802632	250
0.76781	0.879154	0.819718	331
0.769737	0.754839	0.762215	155
0.821561	0.843511	0.832392	262
0.812057	0.78157	0.796522	293
0.735178	0.696629	0.715385	267
0.649402	0.708696	0.677755	230
0.678788	0.643678	0.660767	174
0.733096	0.830645	0.778828	248
0.907216	0.792793	0.846154	222

In [20]:

```
# Logistic Regression model

kfold = KFold(n_splits=10, random_state=7)
model = LogisticRegression()
model.fit(X_train,Y_train)
results = cross_val_score(model, X, Y, cv=kfold)
print('Logistic Regression model Accuracy:\nMean=',round(results.mean()*100,3)
, '%\tStd deviation= ',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actu
al,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : s
upport}
print("\\n",tabulate(results,headers="keys"))
```

Logistic Regression model Accuracy:
Mean= 45.922 % Std deviation= 6.39 %

precision	recall	fscore	support
0.644737	0.588	0.615063	250
0.485944	0.731118	0.583836	331
0	0	0	155
0.362264	0.366412	0.364326	262
0.434615	0.771331	0.555966	293
0.490909	0.202247	0.286472	267
0.342857	0.156522	0.214925	230
0	0	0	174
0.405896	0.721774	0.519594	248
0.770751	0.878378	0.821053	222

In [21]:

```
# Linear Discrimination model

num_folds = 10
kfold = KFold(n_splits=10, random_state=7)
model = LinearDiscriminantAnalysis()
model.fit(X_train,Y_train)
results = cross_val_score(model,X, Y, cv=kfold)
print('Linear Discrimination model Accuracy:\nMean=',round(results.mean()*100,
3), '%\tStd deviation= ',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actu
al,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : s
upport}
print("\\n",tabulate(results,headers="keys"))
```

Linear Discrimination model Accuracy:

Mean= 51.744 % Std deviation= 4.474 %

precision	recall	fscore	support
0.560137	0.652	0.602588	250
0.51	0.616314	0.55814	331
0.230769	0.0967742	0.136364	155
0.505208	0.370229	0.427313	262
0.555838	0.74744	0.637555	293
0.369792	0.265918	0.309368	267
0.456835	0.552174	0.5	230
0.525773	0.293103	0.376384	174
0.600559	0.866935	0.709571	248
0.921212	0.684685	0.78553	222

D:\Anaconda\lib\site-packages\sklearn\discriminant_analysis.py:442
: UserWarning: The priors do not sum to 1. Renormalizing
UserWarning)

D:\Anaconda\lib\site-packages\sklearn\discriminant_analysis.py:442
: UserWarning: The priors do not sum to 1. Renormalizing
UserWarning)

D:\Anaconda\lib\site-packages\sklearn\discriminant_analysis.py:442
: UserWarning: The priors do not sum to 1. Renormalizing
UserWarning)

In [22]:

```
# KNeighborsClassifier model

num_folds = 10
kfold = KFold(n_splits=10, random_state=7)
model = KNeighborsClassifier(n_neighbors=15)
model.fit(X_train,Y_train)
results = cross_val_score(model, X, Y, cv=kfold)
print('KNeighborsClassifier model Accuracy:\nMean=',round(results.mean()*100,3), '%\tStd deviation= ',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actual,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : support}
print("\\n",tabulate(results,headers="keys"))
```

KNeighborsClassifier model Accuracy:

Mean= 63.129 % Std deviation= 3.599 %

precision	recall	fscore	support
0.743961	0.616	0.673961	250
0.582589	0.78852	0.67009	331
0.5	0.303226	0.37751	155
0.694561	0.633588	0.662675	262
0.629179	0.706485	0.665595	293
0.556017	0.501873	0.527559	267
0.454225	0.56087	0.501946	230
0.5	0.33908	0.40411	174
0.623333	0.754032	0.682482	248
0.860465	0.666667	0.751269	222

In [23]:

```
# Bagging Classifier model

seed = 7
kfold = KFold(n_splits=10, random_state=seed)
cart = DecisionTreeClassifier()
num_trees = 130
model = BaggingClassifier(base_estimator=cart, n_estimators=num_trees,max_features = 5, random_state=seed)
model.fit(X_train,Y_train)
results = cross_val_score(model, X, Y, cv=kfold)
print('Bagging Classifier Accuracy:\nMean=',round(results.mean()*100,3), '%\tStd deviation= ',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actual,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : support}
print(" \n",tabulate(results,headers="keys"))
```

Bagging Classifier Accuracy:

Mean= 75.424 % Std deviation= 2.271 %

precision	recall	fscore	support
0.868644	0.82	0.843621	250
0.792507	0.830816	0.811209	331
0.742857	0.670968	0.705085	155
0.8125	0.89313	0.850909	262
0.8157	0.8157	0.8157	293
0.746888	0.674157	0.708661	267
0.663934	0.704348	0.683544	230
0.692308	0.672414	0.682216	174
0.745455	0.826613	0.783939	248
0.894472	0.801802	0.845606	222

In [24]:

```
# Random forest classifier model

num_trees = 150
max_features = 6
kfold = KFold(n_splits=10, random_state=7)
model = RandomForestClassifier(n_estimators=num_trees, max_features=max_features)
model.fit(X_train, Y_train)
results = cross_val_score(model, X, Y, cv=kfold)
print('Random forest classifier Accuracy:\nMean=',round(results.mean()*100,3),
'%\tStd deviation= ',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actual,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : support}
print("\\n",tabulate(results,headers="keys"))
```

Random forest classifier Accuracy:

Mean= 77.731 % Std deviation= 2.442 %

precision	recall	fscore	support
0.869748	0.828	0.848361	250
0.8	0.833837	0.816568	331
0.730496	0.664516	0.695946	155
0.826241	0.889313	0.856618	262
0.844595	0.853242	0.848896	293
0.800885	0.677903	0.73428	267
0.674419	0.756522	0.713115	230
0.691011	0.706897	0.698864	174
0.762846	0.778226	0.770459	248
0.874419	0.846847	0.860412	222

In [25]:

```
# Gradient Boosting Classifier model

seed = 7
num_trees = 100
kfold = KFold(n_splits=10, random_state=seed)
model = GradientBoostingClassifier(n_estimators=num_trees, random_state=seed)
model.fit(X_train,Y_train)
results = cross_val_score(model, X, Y, cv=kfold)
print('Gradient Boosting Classifier Accuracy:\nMean=',round(results.mean()*100
,3), '%\tStd deviation= ',round(results.std()*100,3), '%')

prediction = model.predict(X_test)
precision,recall,fscore,support = metrics.precision_recall_fscore_support(actu
al,prediction,average=None)
results = {"precision":precision,"recall":recall,"fscore":fscore,"support" : s
upport}
print("\\n",tabulate(results,headers="keys"))
```

Gradient Boosting Classifier Accuracy:

Mean= 78.423 % Std deviation= 2.679 %

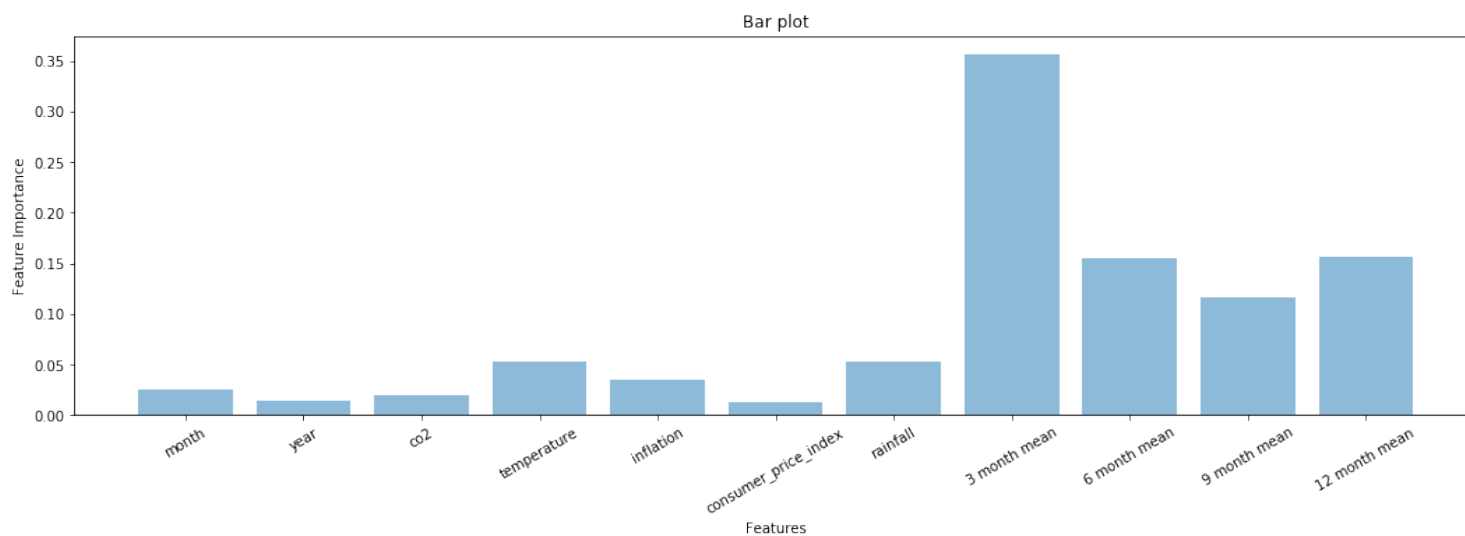
precision	recall	fscore	support
0.882096	0.808	0.843424	250
0.788301	0.854985	0.82029	331
0.774436	0.664516	0.715278	155
0.816901	0.885496	0.849817	262
0.857143	0.83959	0.848276	293
0.784483	0.681648	0.729459	267
0.660079	0.726087	0.691511	230
0.674033	0.701149	0.687324	174
0.742424	0.790323	0.765625	248
0.857143	0.810811	0.833333	222

In [26]:

```
# Feature importances for Gradient Boosting Classifier model
```

```
print('\nFeature importances:',model.feature_importances_)
y_pos = np.arange(cols)
x = model.feature_importances_
f, ax = plt.subplots(figsize=(18,5))
plt.bar(y_pos, x, width=0.8,align='center', alpha=0.5)
plt.title('Bar plot')
plt.xticks(y_pos, df.columns,fontsize=10,rotation=30)
plt.ylabel('Feature Importance')
plt.xlabel('Features')
plt.show()
```

```
Feature importances: [0.026 0.014 0.02  0.054 0.034 0.014 0.053 0.
357 0.155 0.117 0.157]
```



In []: